

# introduction to Bayesian state space models

Ben Bolker

McMaster University  
Departments of Mathematics & Statistics and Biology

20 June 2017

# Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting

# Outline


- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting

# Big picture

- use Bayes' rule
- avoid frequentist contortions
- integrate prior knowledge
- make coherent decisions
- compute hard things

# Bayes rule

$$\underbrace{P(H_i|D)}_{\text{posterior}} = \underbrace{P(D|H_i)}_{\text{likelihood}} \underbrace{P(H_i)}_{\text{prior}} / \underbrace{\sum_j P(H_j)P(D|H_j)}_{P(\text{data})}$$



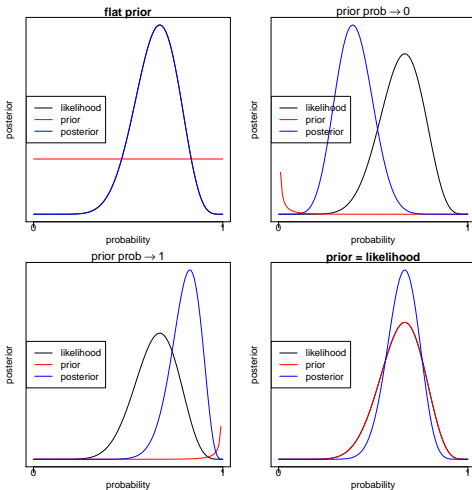
priors:  $P(H_i)$

- usually framed as “prior belief”
- controversial because subjective
- if we set all  $P(H_i)$  equal,  $P(H_i|D) = P(D|H_i) / \sum P(D|H_j)$   
(scaled likelihood)
- can't really be swept under the rug

# more on priors

- **weak** or **diffuse**: little information
- **uninformative** or **flat**: no information
- **improper**: integral diverges (but sometimes OK)
- weak priors can cause problems with sparse data and/or weakly identifiable models
- **conjugate** priors: convenient functional forms

# effects of priors





# Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo**
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting

# Markov chain Monte Carlo

- **general** method for sampling posterior probability densities
- construct a Markov chain whose stationary density equals the desired posterior probability density
- avoids computation of Bayes' rule denominator ( $\iint P(\boldsymbol{\theta}) d\boldsymbol{\theta}$ )

# you won't believe these two MCMC tricks

**Gibbs sampling** sample parameters one at a time,  
exploiting conditioning

**Rejection sampling** (Metropolis-Hastings): pick new values of  
parameters at random, then pick a random number to  
decide whether to keep them

# Gibbs sampling

Because  $\text{Prob}(A|B) = \text{Prob}(A, B)/\text{Prob}(B)$ , we can say

$$\text{Prob}(A, B, C, \dots Z) \propto \text{Prob}(A|B, \dots, Z) \cdot \text{Prob}(B|C, \dots, Z) \cdot \dots \cdot \text{Prob}(Z)$$

This means that we can sample the conditional probabilities  
**sequentially** and get the right answer.

**picture of sampling**

# Metropolis-Hastings

Jump, evaluate (prior  $\times$  likelihood), decide whether to accept

$$\frac{\text{Prob}(A)}{\text{Prob}(B)} = \frac{P(\text{jump } B \rightarrow A) \cdot P(\text{accept } A|B)}{P(\text{jump } A \rightarrow B) \cdot P(\text{accept } B|A)}$$

In the long run our chain will converge to the right distribution

- **candidate distribution**: anything sensible  
(bad choices make sampling slow, but not incorrect)
- **acceptance rule**:

$$P(\text{accept } \theta_2) = \min \left( 1, \frac{\text{Pr}(\theta_2)L(\theta_2)}{\text{Pr}(\theta_1)L(\theta_1)} \right)$$

i.e. “always accept if  $\theta_2$  better: sometimes if  $\theta_2$  worse”

# Magic black boxes

- Can construct your own, customized samplers (Bolker et al., 2003)
- Use BUGS (Bayesian Inference Using Gibbs Sampling) language (WinBUGS, OpenBUGS, JAGS, NIMBLE)
- interfaces from R  
(<https://CRAN.R-project.org/package=R2jags>) or MATLAB  
(<https://github.com/msteyvers/matjags>).

# Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models**
- 4 Other approaches to nonlinear dynamical fitting

# State space models

- address fundamental problem:  $\text{prob}(\text{observations})$  depends on **unobserved** true values
- have to deal with/integrate over all possible values (**latent variables**)
- very high-dimensional: brute force fails
- use the previous two tricks


**important:** often need to handle obs+proc error in some way to get reliable answers and CIs (King et al., 2015)



# BUGS code for the logistic function

```
model <- function() {  
  t[1] <- n0      ## initial values ...  
  o[1] ~ dnorm(t[1],tau.obs)  
  for (i in 2:N) {  ## step through observations ...  
    v[i] <- t[i-1]+r*t[i-1]*(1-t[i-1])/K  
    t[i] ~ dnorm(v[i],tau.proc)  
    o[i] ~ dnorm(t[i],tau.obs)  
  }  
  r ~ dunif(0.1,maxr) ## priors ...  
  ## rate and scale of gamma  
  K ~ dgamma(0.005,0.005)  
  tau.obs ~ dgamma(0.005,0.005)  
  tau.proc ~ dgamma(0.005,0.005)  
  n0 ~ dgamma(1,n0rate)  
}
```

# Dependency structure for logistic model



# BUGS vs R

- BUGS code is not sequential (!)
- BUGS is not vectorized (need for loops)
- BUGS: `~` means “distributed as” (“stochastic node”)  
`<-` means assignment (“logical node”)
- different distribution names and parameterizations (e.g. `dnorm(mean,prec)` for Normal, `dbin(size,prob)` for binomial): see LeBauer et al. (2013)

# Running JAGS

- **Good news:** JAGS code is (relatively) intuitive
- **Bad news:**
  - Debugging is hard
  - Need to figure out how long to run chains (convergence diagnostics)
  - Poor mixing
  - Slow computation

# Running JAGS (details)

- specify model and priors
- get model to compile
- run multiple chains
  - discard **burn-in**
  - thin results
- assess convergence

# Troubleshooting JAGS

- simplify model
- specify initial values explicitly
- narrow priors and/or fix some parameters
- run longer and thin more


# Speeding up MCMC

- reparameterize
- try Hamiltonian MC Carpenter et al. (2016)
- block (correlated) updates to some parameters
- narrow priors and/or fix some parameters

# Diagnostics example

**trace plots:** should look like white noise

```
library(R2jags); library(coda); library(lattice)  
xyplot(as.mcmc(jagsout))
```





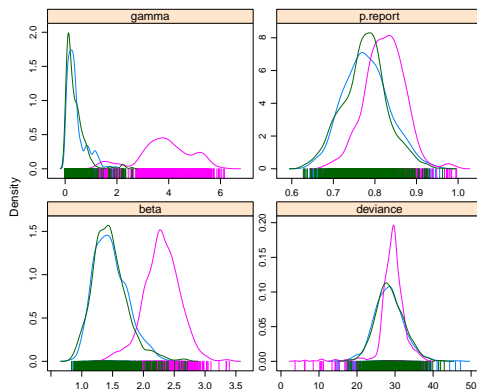
# Diagnostics example

**Gelman-Rubin statistic:** want  $\hat{R} < 1.2$

```
gelman.diag(as.mcmc(jagsout))

## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## beta           2.69         4.78
## deviance        1.01         1.04
## gamma           4.48        13.58
## p.report        1.19         1.54
##
## Multivariate psrf
##
## 3.49
```

# Density plots



# Summary example

```
summary(as.mcmc(jagsout))
```

##	Mean	SD	Naive SE	Time-series SE
## beta	1.7420	0.48411	0.008839	0.023546
## deviance	28.8052	3.73637	0.068217	0.232341
## gamma	1.5430	1.75101	0.031969	0.164691
## p.report	0.7929	0.05669	0.001035	0.003175


# Inference from posteriors

- point estimates: mean, median (marginal)
- interval estimates: quantiles, **highest posterior density** (`coda::HPDinterval`)
- can summarize any quantity computed from samples (e.g. predictions)

## Further resources

- Gelman et al. (2013) (fairly hard-core)
- Gelman and Hill (2006) (more regression-focused)
- Hobbs and Hooten (2015) (friendlier, ecologist-focused)
- McCarthy (2007) (even friendlier)

# The problem with MCMC (xkcd)



# Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting**

# advances in MCMC

- new toolboxes: NIMBLE Li et al. (2017); de Valpine et al. (2017), PyMC
- Hamiltonian Monte Carlo: Stan Carpenter et al. (2016)
- variational Bayes, expectation-propagation . . . Gelman et al. (2013)



# Frequentist alternatives

- MCMC is usually Bayesian;  
opens various cans of worms
- there are many other related approaches, some classical
  - expectation-maximization
  - sequential Monte Carlo/particle filters (Ionides et al., 2006; Doucet et al., 2001; de Valpine, 2004): R pomp, NIMBLE packages, PyMC
  - data cloning (Lele et al., 2007): R dc1one package

# Estimation for continuous-time models

## Particle methods (freq or Bayesian)

- simulate many trajectories (“particles”) step-by-step
- at each step, resample particles weighted by likelihood of current location

## Approximate Bayesian computation

- sample parameters from prior
- simulate trajectories for each parameter set
- compute summary statistics (“probes”)
- save trajectories with summary stats near observed values

# References

- Bolker, B., Okuyama, T., et al., 2003. *Ecological Applications*, 13(3):763–775.
- Carpenter, B., Gelman, A., et al., 2016. *Journal of Statistical Software*, 20.
- de Valpine, P., 2004. *Journal of the American Statistical Association*, 99:523–536.
- de Valpine, P., Turek, D., et al., 2017. *Journal of Computational and Graphical Statistics*, 26(2):403–413. ISSN 1061-8600. doi:10.1080/10618600.2016.1172487.
- Doucet, A., de Freitas, N., and Gordon, N.J., 2001. *Sequential Monte Carlo methods in practice*. Springer-Verlag, New York, USA.
- Gelman, A., Carlin, J.B., et al., 2013. *Bayesian Data Analysis*. Chapman and Hall/CRC, Boca Raton, 3d edition. ISBN 978-1-4398-4095-5.
- Gelman, A. and Hill, J., 2006. *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge University Press, Cambridge, England.
- Hobbs, N.T. and Hooten, M.B., 2015. *Bayesian Models: A Statistical Primer for Ecologists*. Princeton University Press, Princeton, New Jersey. ISBN 978-0-691-15928-7.
- Ionides, E.L., Bretó, C., and King, A.A., 2006. *Proceedings of the National Academy of Sciences of the USA*, 103(49):18438–18443. doi:doi:10.1073/pnas.0603181103.
- King, A.A., Cellès, M.D.d., et al., 2015. *Proc. R. Soc. B*, 282(1806):20150347. ISSN 0962-8452, 1471-2954. doi:10.1098/rspb.2015.0347.
- LeBauer, D.S., Dietze, M.C., and Bolker, B.M., 2013. *R Journal*, 5(1):207–209.
- Lele, S.R., Dennis, B., and Lutscher, F., 2007. *Ecology Letters*, 10:551–563. doi:doi:10.1111/j.1461-0248.2007.01047.x.
- Li, M., Dushoff, J., and Bolker, B.M., 2017. *bioRxiv*, page 110767.
- McCarthy, M., 2007. *Bayesian methods for ecology*. Cambridge University Press, Cambridge, England.