

State space models

Ben Bolker

McMaster University
Departments of Mathematics & Statistics and Biology

20 June 2017

Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting

Outline


- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting

Big picture

- use Bayes' rule
- avoid frequentist contortions
- integrate prior knowledge
- make coherent decisions
- compute hard things

Bayes rule

$$\underbrace{P(H_i|D)}_{\text{posterior}} = \underbrace{P(D|H_i)}_{\text{likelihood}} \underbrace{P(H_i)}_{\text{prior}} / \underbrace{\sum_j P(H_j)P(D|H_j)}_{P(\text{data})}$$



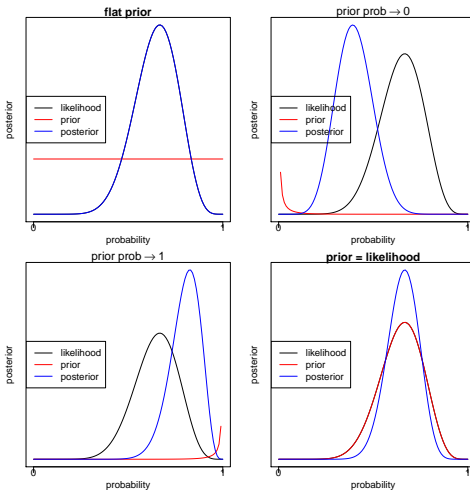
priors: $P(H_i)$

- usually framed as “prior belief”
- controversial because subjective
- if we set all $P(H_i)$ equal, $P(H_i|D) = P(D|H_i) / \sum P(D|H_j)$
(scaled likelihood)
- can't really be swept under the rug

more on priors

- **weak** or **diffuse**: little information
- **uninformative** or **flat**: no information
- **improper**: integral diverges (but sometimes OK)
- weak priors can cause problems with sparse data and/or weakly identifiable models
- **conjugate** priors: convenient functional forms

effects of priors



Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo**
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting

Markov chain Monte Carlo

- **general** method for sampling posterior probability densities
- construct a Markov chain whose stationary density equals the desired posterior probability density
- avoids computation of Bayes' rule denominator ($\iint P(\boldsymbol{\theta}) d\boldsymbol{\theta}$)

you won't believe these two MCMC tricks

Gibbs sampling exploit conditioning, sample parameters one at a time

Rejection sampling (Metropolis-Hastings): pick new values of parameters at random, then pick a random number to decide whether to keep them

Gibbs sampling

Because $\text{Prob}(A|B) = \text{Prob}(A, B)/\text{Prob}(B)$, we can say

$$\text{Prob}(A, B, C, \dots Z) \propto \text{Prob}(A|B, \dots, Z) \cdot \text{Prob}(B|C, \dots, Z) \cdot \dots \cdot \text{Prob}(Z)$$

This means that we can sample the conditional probabilities
sequentially and get the right answer.

picture of sampling

Metropolis-Hastings

Jump, evaluate (prior \times likelihood), decide whether to accept

$$\frac{\text{Prob}(A)}{\text{Prob}(B)} = \frac{P(\text{jump } B \rightarrow A)P(\text{accept } A|B)}{P(\text{jump } A \rightarrow B)P(\text{accept } B|A)}$$

In the long run our chain will converge to the right distribution

- **candidate distribution**: anything sensible
- **acceptance rule**:

$$P(\text{accept } \theta_2) = \min \left(1, \frac{\text{Pr}(\theta_2)L(\theta_2)}{\text{Pr}(\theta_1)L(\theta_1)} \right)$$

i.e. “always accept if θ_2 better: sometimes if θ_2 worse”

Magic black boxes

- Can construct your own, customized samplers (?)
- Use BUGS (Bayesian Inference Using Gibbs Sampling) language (WinBUGS, OpenBUGS, JAGS, NIMBLE)
- interfaces from R (R2jags package) or MATLAB (<https://github.com/msteyvers/matjags>).

Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models**
- 4 Other approaches to nonlinear dynamical fitting


State space models

- address fundamental problem: $\text{prob}(\text{observations})$ depends on **unobserved** true values
- have to deal with/integrate over all possible values (**latent variables**)
- very high-dimensional: brute force fails
- use the previous two tricks

BUGS code for the logistic function

```
model <- function() {  
  t[1] <- n0      ## initial values ...  
  o[1] ~ dnorm(t[1],tau.obs)  
  for (i in 2:N) {  ## step through observations ...  
    v[i] <- t[i-1]+r*t[i-1]*(1-t[i-1])/K  
    t[i] ~ dnorm(v[i],tau.proc)  
    o[i] ~ dnorm(t[i],tau.obs)  
  }  
  r ~ dunif(0.1,maxr) ## priors ...  
  ## rate and scale of gamma  
  K ~ dgamma(0.005,0.005)  
  tau.obs ~ dgamma(0.005,0.005)  
  tau.proc ~ dgamma(0.005,0.005)  
  n0 ~ dgamma(1,n0rate)  
}
```

Dependency structure for logistic model



BUGS vs R

- BUGS code is not sequential (!)
- BUGS is not vectorized (need for loops)
- BUGS: `~` is “distributed as” (“stochastic node”), `<-` is assignment (“logical node”)
- different distribution names and parameterizations (e.g. `dnorm(mean,prec)` for Normal, `dbin(size,prob)` for binomial): see LeBauer et al. (2013)

Running JAGS

- **Good news:** JAGS code is (relatively) intuitive
- **Bad news:**
 - Debugging is hard
 - Need to figure out how long to run chains (convergence diagnostics)
 - Poor mixing
 - Slow computation

Running JAGS (details)

- specify model and priors
- get model to compile
- run multiple chains
- assess convergence

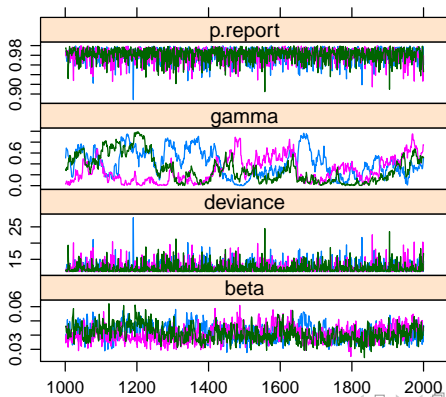
Troubleshooting JAGS

- simplify model
- specify initial values explicitly
- narrow priors and/or fix some parameters
- run longer and thin

Diagnostics example

trace plots: should look like white noise

```
library(R2jags); library(coda); library(lattice)
xyplot(as.mcmc(jagsout))
```



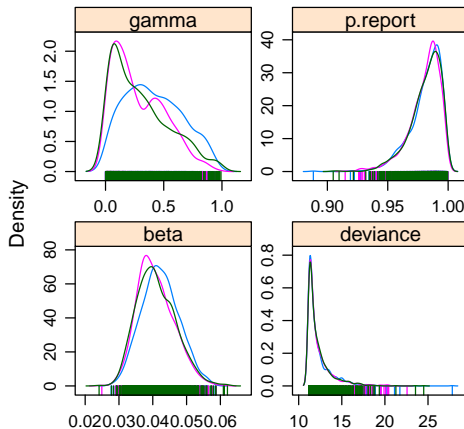
Diagnostics example

Gelman-Rubin statistic: want $\hat{R} < 1.2$

```
gelman.diag(as.mcmc(jagsout))

## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## beta           1.01         1.05
## deviance        1.00         1.00
## gamma           1.06         1.20
## p.report        1.00         1.00
##
## Multivariate psrf
##
## 1.05
```


Density plots




Summary example

```
summary(as.mcmc(jagsout))
```

##		Mean	SD	Naive SE	Time-series SE
## beta		0.04138	0.005566	0.0001016	0.0003900
## deviance		12.27502	1.506429	0.0275035	0.0290581
## gamma		0.33621	0.243174	0.0044397	0.0452345
## p.report		0.98090	0.013220	0.0002414	0.0002512

The problem with MCMC (xkcd)



Outline

- 1 super-quick intro to Bayes
- 2 Markov chain Monte Carlo
- 3 State-space models
- 4 Other approaches to nonlinear dynamical fitting**

Frequentist methods

- MCMC is usually done in a Bayesian framework;
opens various cans of worms
- there are many other related approaches, some classical
 - sequential Monte Carlo/particle filters (Ionides et al., 2006; Doucet et al., 2001; de Valpine, 2004): R pomp package
 - data cloning (Lele et al., 2007): R dc1one package

Continuous-time models

- Particle methods
- Approximate Bayesian computation

References

- de Valpine, P., 2004. *Journal of the American Statistical Association*, 99:523–536.
- Doucet, A., de Freitas, N., and Gordon, N.J., 2001. *Sequential Monte Carlo methods in practice*. Springer-Verlag, New York, USA.
- Ionides, E.L., Bretó, C., and King, A.A., 2006. *Proceedings of the National Academy of Sciences of the USA*, 103(49):18438–18443. doi:doi:10.1073pnas.0603181103.
- LeBauer, D.S., Dietze, M.C., and Bolker, B.M., 2013. *R Journal*, 5(1):207–209.
- Lele, S.R., Dennis, B., and Lutscher, F., 2007. *Ecology Letters*, 10:551–563. doi:doi:10.1111/j.1461-0248.2007.01047.x.